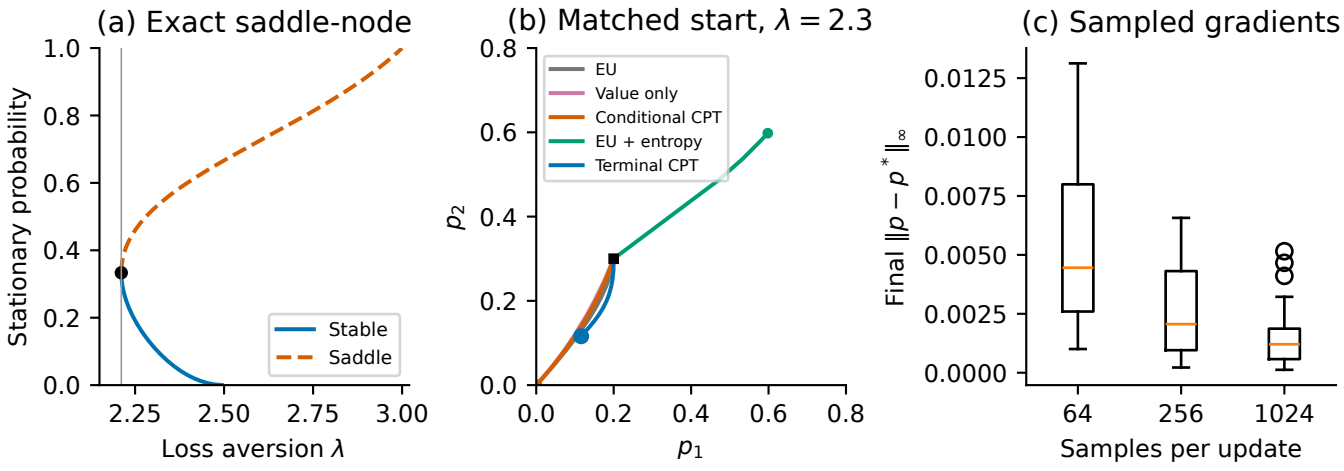
(a) Exact saddle-node
Stationary probability
Loss aversion λ
Stable
Saddle
(b) Matched start, $\lambda = 2.3$
p_2
p_1
EU
Value only
Conditional CPT
EU + entropy
Terminal CPT
(c) Sampled gradients
Final $\|p - p^*\|_\infty$
Samples per update
64
256
1024